

Scheduling Algorithms

FCFS

- Process added first in ready queue should be scheduled first.
- Non-preemptive scheduling
- Scheduler is invoked when process is terminated, blocked or gives up CPU is ready for execution.
- Convoy Effect: Larger processes slow down execution of other processes.

SJF

- Process with lowest burst time is scheduled first.
- Non-preemptive scheduling
- Minimum waiting time

SRTF - Shortest Remaining Time First

- Similar to SJF - but Preemptive scheduling
- Minimum waiting time

Priority

- Each process is associated with some priority level. Usually lower the number, higher is the priority.
- Preemptive scheduling or Non Preemptive scheduling
- Starvation
 - Problem may arise in priority scheduling.
 - Process not getting CPU time due to other high priority processes.
 - Process is in ready state (ready queue).
 - May be handled with aging -- dynamically increasing priority of the process.

Round-Robin

- Preemptive scheduling
- Process is assigned a time quantum/slice.
- Once time slice is completed/expired, then process is forcibly preempted and other process is scheduled.
- Min response time.

Fair-share

- CPU time is divided into epoch times.
- Each ready process gets some time share in each epoch time.
- Process is assigned a time share in proportion with its priority.
- In Linux, processes with time-sharing (TS) class have nice value. Range of nice value is -20 (highest priority) to +19 (lowest priority).

IPC overview

- A process cannot access the memory of another process directly. OS provides IPC mechanisms so that processes can communicate with each other.
- IPC models
 - Shared memory model
 - Processes write/read from the memory region accessible to both the processes.
 - OS only provides access to the shared memory region.
 - Message passing model
 - Process send message to the OS and the other process receives message from the OS.
 - This is slower compared to shared memory model.
- Unix/Linux IPC mechanisms
 - Signals
 - Shared memory
 - Message queue
 - Pipe
 - Socket

Message Queue

- Used to transfer packets of data from one process to another.
- It is bi-directional IPC mechanism.
- Internally OS maintains list of messages called as "message queue" or "mailbox".
- The info about msg que is stored in a object. It contains unique KEY, permissions, message list, number of messages in list, processes waiting for a message to receive (waiting queue).

Message Queue Syscalls

- `msgget()`
 - Create message queue object.
 - `mqid = msgget(mq_key, flags);`
 - `arg1`: unique key
 - `arg2`: `IPC_CREAT | 0600` to create new message queue.
 - returns message queue id on success
- `msgctl()`
 - Get info about message queue or destroy message queue
 - `msgctl(mqid, IPC_RMID, NULL) --` to destroy message queue
 - `arg1`: message queue id
 - `arg2`: `commad = IPC_RMID` to destroy message queue
 - `arg3`: `NULL` (not required while destroying message queue)
- `msgsnd()` - send message into que
 - Send message in the message queue.
 - `ret = msgsnd(mqid, msg_addr, msg_size, flags)`
 - `arg1`: message queue id
 - `arg2`: base address of message object
 - `arg3`: message body/payload size
 - `arg4`: flags (=0 for default behaviour)
 - returns 0 on success.
- `msgrcv()` - receive message from que

- Receive message from the message queue of given type.
- `ret = msgrcv(mqid, msg_addr, msg_size, msg_type, flags)`
 - `arg1`: message queue id
 - `arg2`: base address of message object to collect message (out param)
 - `arg3`: message body/payload size
 - `arg4`: type of message to be received
 - `arg5`: flags (=0 for default behaviour)
 - returns number of bytes (body) received on success.

Signals

- OS have a set of predefined signals, which can be displayed using command
 - `terminal> kill -l`
- A process can send signal to another process or OS can send signal to any process.
- Information about signals.
 - `terminal> man 7 signal`

Send signal

- `kill` command is used to send signal to another process, which internally use `kill()` syscall.
 - `terminal> kill -SIG pid`
- `pkill` command is used to send signal to multiple processes/instances of the same program.
 - `terminal> pkill -SIG programname`

Imporant Signals

1. **SIGINT (2)**: When CTRL+C is pressed, INT signal is sent to the foreground process.
2. **SIGTERM (15)**: During system shutdown, OS send this signal to all processes. Process can handle this signal to close resources and get terminated.
3. **SIGKILL (9)**: During system shutdown, OS send this signal to all processes to forcefully kill them. Process cannot handle this signal.
4. **SIGSTOP (19)**: Pressing CTRL+S, generate this signal which suspend the foreground process. Process cannot handle this signal.
5. **SIGCONT (18)**: Pressing CTRL+Q, generate this signal which resume suspended the process.
6. **SIGSEGV (11)**: If process access invalid memory address (dangling pointer), OS send this signal to process causing process to get terminated. It prints error message "Segmentation Fault".
7. **SIGCHLD (17)**: When child process is terminated, this signal is sent to the parent process. The parent process may handle this to get the exit code of the child (`wait()` syscall).
8. **SIGHUP (1)**: When a terminal is closed, all processes running in that terminal and terminated due to Hang up signal.

Signals related syscalls

- `kill()` syscall
 - `kill()` send signal to another proces.
 - `ret = kill(pid, signum);`
 - `arg1`: pid of the process to whom signal is to be sent.
 - `arg2`: signal number -- defined in `signal.h`

- returns: 0 on success and -1 on failure.
- `signal() syscall`
 - `signal()` is used to install signal handler in current process (signal handler table).
 - When signal is received, OS calls registered signal handler.
 - `old_handler = signal(signum, new_handler);`
 - `arg1 (int)`: signal number of signal to be handled (except SIGKILL and SIGSTOP).
 - `arg2 (fn ptr)`: address of signal handler function.
 - `typedef void (*sighandler_t)(int);`
 - returns: address of old signal handler (in the table).

File IO syscalls

- `open()`
- `read()`
- `write()`
- `close()`
- `lseek()`

`open() syscall`

- `fd = open("file-path", flags, mode);`
 - `arg1`: path of file to be opened
 - `arg2`: flags - how you want to open the file
 - `O_RDONLY, O_WRONLY, O_RDWR` -- read-write flags
 - `O_TRUNC` -- delete the contents of file while opening
 - `O_APPEND` -- write at the end of file
 - `O_CREAT` -- create a new file (if not present) -- must give `arg3`
 - `arg3`: mode - permissions for new file - octal number
 - returns file descriptor on success, and -1 on failure.
 - file descriptor is int that uniquely identifies the file in that process.
 - `fd` is used in other file io syscalls e.g. `close()`, `read()`, `write()`, `lseek()`.

`close() syscall`

- `int close(fd)`
 - `arg1`: fd returned by `open()`.
 - Returns 0 on success and -1 on error
- Decrement ref count in open file table entry (struct file).
- If ref count drops to zero, OFT entry is deleted (from OFT).

`read() syscall`

- `count = read(fd, buf, length);`
 - `arg1`: file descriptor
 - `arg2`: buffer in which you want to read
 - `arg3`: length of buffer
 - Returns: Number of characters read

write() syscall

- `ssize_t write(int fd, const void *buf, size_t count);`
 - arg1: file descriptor
 - arg2: buffer from which you want to write
 - arg3: length of data
 - Returns: Number of characters written

lseek() syscall

- Change the file position in open file table entry (struct file --> f_pos).
- And returns new file position (from the beginning of the file).
- `off_t lseek(int fd, off_t offset, int whence);`
- Examples:
 - `lseek(fd, 0, SEEK_SET);`
 - `filp->f_pos = 0;`
 - `lseek(fd, offset, SEEK_SET);`
 - `filp->f_pos = offset;`
 - `lseek(fd, 0, SEEK_END);`
 - `filp->f_pos = size; // file size (from the inode)`
 - `lseek(fd, offset, SEEK_END);`
 - `filp->f_pos = size + offset; // note that offset will be negative`
 - `lseek(fd, offset, SEEK_CUR);`
 - `filp->f_pos = filp->f_pos + offset;`

Pipe

- Pipe is used to communicate between two processes.
- It is stream based uni-directional communication.
- Pipe is internally implemented as a kernel buffer, in which data can be written/read.
- If pipe (buffer) is empty, reading process will be blocked.
- If pipe (buffer) is full, writing process will be blocked.
- If writer process is terminated, reader process will read the data from pipe buffer and then will get EOF.
- If reading process is terminated, writing process will receive SIGPIPE signal.
- There are two types of pipe:
 - Unnamed Pipe
 - Named Pipe

pipe() syscall

- To create unnamed pipe.
- `ret = pipe(fd); // int fd[2];`
 - arg1: array of two ints to collect fd (out param).
 - returns 0 on success.
 - `fd[]` gets two ends of pipe.
 - `fd[0]` -- read end (file descriptor) of pipe
 - `fd[1]` -- write end (file descriptor) of the file